

"A Tale of Two Provers":

A Comparison of Dependent Haskell and F*

Xinyue Zhang

Advised by Richard A. Eisenberg

March, 2018

Abstract

To be filled

Submitted in Partial fulfillment of the requirement for BA in Computer Science at Bryn Mawr College

Contents

1 Introduction and Motivation	4
2 Background	5
2.1 From Imperative to functional programming	5
2.2 Program Specification with Dependent and Refinement Types	7
2.3 Programming Effects	7
2.3.1 Monads	7
2.4 Dependent Haskell	8
2.4.1 Introduction to Haskell	8
2.4.2 Dependently-Typed Dependent Haskell	9
2.5 Program Verification in Dependent Haskell	10
2.5.1 Algebraic Data Types	10
2.5.2 Generalized Algebraic Data Types	11
2.5.3 Type Class	11
2.5.4 A Simple Verification Example	11
2.5.5 Type Family	12
2.5.6 Program Verification with Dependent Types and Type Families	12
2.6 FStar	13
2.6.1 Introduction to F*	13
2.6.1.1 F* Basics	13
2.6.1.2 Simple Inductive Types	14
2.6.2 Intrinsic and Extrinsic Style of Verification	15
2.6.2.1 Proving List Concatenation Intrinsically through Type Enrichment	15
2.6.2.2 Proving List Concatenation Externally through Lemmas	15
2.6.3 Proof Automation	16
2.6.4 Categorizing effects in F*	16
2.6.4.1 Dijkstra Monad and Hoare Logic	16
2.6.4.2 Proving Termination	17
3 Approach and Novelty	18
4 Current Results	19
4.1 User Interface: General Language Design and Support	19
4.1.1 Syntactic Verbosity	19
4.1.1.1 Dependent Haskell	19
4.1.1.2 F*	21
4.1.2 Library and Documentation Supports	22
4.1.3 Proof Complexity	22
4.2 Approaches to Potentially Non-terminating Functions	23
4.2.1 Dependent Haskell might be Complete towards Potentially Divergent Functions	23
4.2.2 F* is Incomplete towards Potentially Divergent Functions	25
5 Related Work and Contributions	28
5.1 Language Comparison	28

5.2	Dependent Haskell	28
5.2.1	Functional Dependency	28
5.2.2	Constrained Type Families	28
5.2.3	Promoting Functions to Type Families in Haskell	28
5.3	F*	29
5.3.1	Relative Completeness in Hoare Logic	29
5.3.2	Dijkstra Monads	29
5.3.3	Heredity Substitutions	29
5.4	The Zombie Language	29
Appendices		29
.1	Insersion Sort	29
.1.1	Dependent Haskell	29
.1.2	F*	29
.2	Peano Division	29
.2.1	Dependent Haskell	29
.2.2	F*	31

1 Introduction and Motivation

"Testing can show the presence of errors, but not their absence."

– Edsger Wybe Dijkstra

Programming languages are powerful tools used to design and build complex computer software in order to perform various tasks. Yet complex software programs often have ~~errors commonly known as~~ bugs that prevent them from functioning as expected. Almost everything these days is dependent on software, from smart devices to vehicles, and to national defense systems. Most of the time bugs cause small problems such as an application crash. However, in many cases, tiny errors in a program cause catastrophic if not life threatening consequences such as a financial crisis or a plane crash. Therefore, various research must be done to ensure the correctness of a program.

Through systematic testing and static checking, it's possible to find some incorrect program behaviors, assuming that it is easy or even possible to check the correctness of the output from the program. Yet exhaustive testing like these can only apply to very few trivial programs, and is both labor intensive and program specific, making it not feasible. As Edsger Wybe Dijkstra stated, "testing can show the presence of errors, but not their absence" ~~is~~ testing cannot guarantee the correctness of a program. Other options for error checking exist, and formal verification is the most widely used technique. Formal verification uses mathematics to ensure the correctness of software with respect to a defined or inferred specification. We say a software is correct if it meets all of its specifications, which define all its intended behaviors. Formal verification can be applied to programming languages to help verify that there are no bugs in the software written. Using formal verification, we want to create a programming language that makes sure all programs written indeed perform as intended. We believe that verification-oriented programming languages have the potential to detect most of the bugs in a program at compile time. Therefore, we are looking for better techniques to support program verification in programming languages.

One major approach in formal verification is type system design. A Type system is a mechanism that enforces rules on programs through defining type constraints and the associated type applications. A type is an abstract value that classifies a group of values sharing some properties. For example, we say that `true` and `false` are both of type `Boolean`. Since the `+` operation is not valid for booleans, we can catch invalid operations such as `true + false` using typing rules. Through designing a formal type system, program specifications can be encoded in the types, allowing the program to become what's known as a proof-carrying code [11].

There are three main ways to specify a program, through dependent type, refinement type and Hoare logic (dependent type and refinement type are discussed in detail in section 2.2 and Hoare logic in section 2.6.4). A dependent type is a type whose definition is predicated upon a dynamic value. As such, dependently-typed programming languages enable programmers to express more detailed specifications with types and to perform more powerful verifications through type checking. On the other hand, refinement type is a type that satisfies a given logical predicates. It helps reduce code redundancy and makes specifications more intuitive. For example, consider a length indexed list `l` and let `n` represent the length of the list and `a` represent the type of the list elements. Then `l` is specified through dependent type as `l:Vec a n` with the user-defined length-indexed `Vec` datatype, and through refinement type as `l>List a{len l = n}` with the user-

provided `len` function. Finally, Hoare logic is a system of logical rules for reasoning about programs. It consists of Hoare triples that specify the precondition, command of execution, and postcondition such that if the precondition holds before running the command, then the postcondition will hold as long as the command terminates.

Programming languages like Coq [2], Agda [12], and Dafny [9], commonly known as theorem provers, have been studied for a long time and are relatively well understood. They laid down a solid foundation for program verification, but are limited when applied to general-purpose applications in real life.^{no. they are proof assistants} However there are various industrial-strength functional programming languages with resemblance to mathematical functions that also provide a good basis for supporting verifications. Researchers on functional programming languages, like Haskell and ~~MetaLanguage~~(ML), gradually start to support verifications with dependent types. Dependent Haskell and F* are two such dependently-typed languages. Each presents a unique and effective program verification technique.

Dependent Haskell and F* are similar in that they both aim towards a verification-based yet general-purpose programming language. They are, however, noticeably different in their unique approaches to verify potential divergent functions. Dependent Haskell doesn't require termination checking, whereas F* categorizes divergent functions into a special DIV effect (See section 2.5 and 2.6.4 for details). Most dependently typed verifiers require compile-time termination checking to ensure the correctness of proofs. However, as Alan Turing proved ^{using} in the Halting Problem, given an arbitrary program and its inputs, checking whether this program will terminate or loop forever is an undecidable problem. Yet functions that diverge or hard to prove termination are very common in practical programming languages. Therefore, we are interested in exploring how to better express and reason about them.

In this undergraduate thesis, we aim to introduce ~~the~~ two verification-oriented programming languages, Dependent Haskell and F*, present their unique type system designs with concrete examples, and give a detailed comparison from a combination of theoretical and practical standpoints. We expect our research to setup a solid foundation for potential collaborations across the two communities, and to possibly propose promising future research directions in formal verifications for the general programming languages community.

2 Background

2.1 From Imperative to functional programming

Most programmers come from a Java or C/C++ background [1]. These languages are designed to primarily support imperative or procedural programming. Imperative style of programming, a developer writes code that details each step for computers to execute. These imperative programs often have side effects that modify program inputs or global variables. In the field of programming languages, any function modifying global states in execution is known as stateful. For example, global variables can be modified in any function in scope. Functional programming, as its name suggests, models a problem as a set of functions to be executed, detailing the inputs and outputs for each function. Some common functional programming languages are Haskell, ML, Scala, and Lisp. Similar to the mathematical functions introduced in high school algebra, ~~one~~^{one of these languages} abstracts the idea of function and allow it to operate on any structured data types, such as boolean, string, list, etc. Functional programs are usually executed by evaluating expressions and typically

|| awk

Second Draft

avoid mutating program states. In functional programming, functions are treated like any other values and can be passed as arguments to other functions or be returned as a result of a function. Functions that take functional arguments are called higher-order functions.

One major difference between the imperative and functional style of programming is how they handle program iterations. Imperative programming languages usually support both loops and recursion, however functional programming languages depend highly on recursion and higher-order functions, like `map`, `filter` etc. For example, if one wants to double every element of the list `[1,5,2,4,3]`, one can implement in Java, an imperative programming language as shown below

```
public void double(int[] arr) {  
    for(int i = 0; i < arr.length; i++) {  
        arr[i] = 2 * arr[i];  
    }  
}
```

That's very subjective. Understanding that loop requires a lot of knowledge about loops & arrays.

The above code is very intuitive, but notice that the input list `[1,5,2,4,3]` is directly modified after the execution of the function, `double`, to the result `[2,10,4,8,6]`.

The same operation can be defined in Haskell in functional style with the higher-order function `map` as seen in the code below,

```
double :: [Int] -> [Int]  
double = map (*2)
```

*NR: not all functional langs
are typed.*

*too much new in
this example:
- higher-order function
- Haskell sections
- partial application
- currying.*

*better
to break
this into
multiple
examples*

In functional programming style, programmers often declare type signatures of a function on the first line and give actual implementations of it on the following lines. In the `double` example above, the input is specified as `[Int]`, denoting a list of integers, and output also as `[Int]`. The entire type definition `[Int] -> [Int]` after the double colon belongs to the type level. Specifically, when applying `double` to the input list as in `double [1,5,2,4,3]`, the function `(*2)` is mapped to every element of the input list, returning the doubled list `[2,10,4,8,6]` as a result without modifying the input list. Instead of using parentheses to symbolize function application, one can simply put a space between function name and its parameters. The function, `(*2)`, applied here uses the idea of currying. The expression `(*2)` is equivalent to the familiar lambda notation `\x -> x * 2`, as both functions take a number and double it. The idea of currying is widely used in functionally styled programming, since functions with too few parameters will simply return another function that takes as many parameters as left out. Here, `*` has type int \rightarrow `int -> int`, i.e. taking two integer inputs and returns an integer output as the product. As only one parameter is applied to `*` when calling `(*2)`, this expression essentially becomes another function that takes in the remaining integer parameter and returns the product of that integer and 2. Finally, applying the higher-order function `map`, the doubling operation is mapped on every element of the list, doubling the entire list.

*This is too
terse to be
helpful.*

2.2 Program Specification with Dependent and Refinement Types

As discussed in the introduction section, the starting point of any program verification is to specify its properties. Programs can be specified using either dependent types or refinement types. To illustrate these definitions in more detail, consider the example of defining a list. In imperative programming languages like Java or C/C++, elements in a list must all share a single type, such as the boolean type. Take Java as an example, a ~~list~~^{array} of length 5 will be ~~declared as boolean[5]~~. To specify the list carrying its own length in verification-oriented programming languages, one can predicate the type of this list of booleans v on its length 5, as in the type $\text{Vec } 5 \text{ Bool}$, with the defined length-indexed Vec datatype. v 's type $\text{Vec } 5 \text{ Bool}$ is said to be dependently-typed as it's indexed on the number 5 of another type Nat (the type Nat represents the natural numbers, defined to be non-negative integers). Similarly, this list of booleans of length 5 can be specified through logical predicates using refinement types. Recall that refinement types specify a logical formula that the types have to additionally satisfy. So one must first define a function len calculating the length of a list and then apply it in the predicate to produce the refinement type $v : \text{List Bool} \{ \text{len } v = 5 \}$. The logical formula $\text{len } v = 5$ is called the ~~subtyping~~^{subsumption} predicate of the refinement type, and List is the built-in data type from Haskell's `Data.List` library. Both dependent type and refinement type are commonly applied in program verifications. Dependent Haskell only supports dependent types, but F* supports both dependent types and refinement types. We will compare them in detail in our result section.

I'm lost here; what List comes from Data.List?
Haskell doesn't even have refinement types.

2.3 Programming Effects

Most verification techniques nowadays are based on pure evaluations, receiving inputs and directly returning outputs without modifying any other program states. However, industrial-strength programming in the real world often has side effects. A practical program usually doesn't just return a value output purely from operating on its inputs. Instead, the program might modify other states beyond its scope, interact with its calling function or the outside world through streams and sockets, or run forever without termination.

what's your basis for this claim? Sure it can do many things exactly this.

2.3.1 Monads

In most functional programming languages, one common way to express programming effects is by using monads. Through monads, programmers can structure computations in terms of values and sequences of computations using these values. Without interacting with the external states, monads allow effectful computations to be embedded in pure languages.

no, just a few languages. These sentences don't say much.

The term monad might be overwhelming for programmers new to Haskell, but it's actually a simple idea. A monad is essentially wrapping a value into the wrapped value. Additionally, monads support unwrapping the wrapped value to reveal the enclosed value. Suppose that one wants to write a function to calculate the square root of an integer and to return the result in double. One attempt is to specify the function with type $\text{sqrt} : \text{Int} \rightarrow \text{Double}$. Yet one quickly finds oneself in trouble as negative integers don't have real square roots. In languages like Java or C/C++, one can resolve the complexity by simply returning `null` for any negative input. However, functions in Haskell must always return

but double in Java doesn't include null

values of the same type, so the untyped value `null` is not supported in Haskell. Instead, Haskell introduce the `Maybe` monad, defined as

```
data Maybe a = Just a | Nothing ← your readers will  
not understand this syntax
```

The `Maybe` monad could either represents a wrapped value incorporating a successfully evaluated value or denoting the state of potential failure. but the fact that `Maybe` is a monad is irrelevant in this example.

In the `sqrt` example, successful evaluation of the nonnegative square roots can be denoted as `Just result`, where `result` is the resulting square root. Whereas the invalid operation on negative integers can return the value `Nothing`, representing the evaluation failure. The `Nothing` value can be seen as the analogue to the special `null` term.

Beyond the `Maybe` monad, Haskell defines the `IO` monad for handling inputs and outputs, `Error` monad for exception handling, `State` monad to simulate in-place value updates etc.

*This section tells me that you
do not really understand monads — it might be helpful
to look around online
for explanations of
the concept*

*this \$ will
not help
your reader.*

2.4 Dependent Haskell

2.4.1 Introduction to Haskell

Haskell is a statically typed and strongly typed purely functional programming language. It features a type system with type inference and lazy evaluation. ← not a type system feature

As a pure functional programming language, each function in Haskell only takes in some inputs and returns an output after some operations, without modifying any unspecified program states. Unlike many functions (which should really be called procedures) in an impure language, functions in Haskell can only observe the inputs defined but not modifying them. Haskell functions don't mutate variables or handle errors. They also don't read or write to standard inputs or outputs, nor do they connect to sockets to communicate with the outside world. This seemingly limiting feature ensures that each execution of the function always returns exactly the same result, which allows formal proofs on functions and also composing functions to perform more complex operations.

Haskell is also statically typed, where every variable's type is determined at compile time. In addition, it has a type system that ensures every variable type to be checked before performing any operations. It also enforces strict typing rules on functions, also known as strongly-typed. In strongly-typed languages, every variable type is predefined and requires explicit casts when used differently. By this definition, Java is strongly typed since it rejects operations like `1 + "Haskell"` but C is weakly typed since it accepts the mixed-typed operations like `1 + false`, treating `false` as 0.

Even though every single type is strictly enforced, Haskell still remains elegant and concise. It uses a type system with full support of type inference, enabling programmers to leave out many type definitions. For example, code snippets like `a = x + 2` that adds an integer variable `x` to a numeric literal 2 is allowed without explicitly specifying 2's type. Haskell can deduce 2 to have type `Int` from the operation `(+)`, since `(+)` supports only addition of numbers of the same type. but you haven't said what type x has!

*This is what
a type system
always does*

*bad example —
requires knowing
about Haskell's
overloaded
types*

Finally, Haskell is lazy, since every Haskell operation is lazily evaluated. This means that the final result of an operation is not computed unless they're required. Suppose we want to quadruple the list `l = [1,2,3,4]` using our afore-defined `double` function by calling `r = double(double [1,2,3,4])`. In non-lazy languages, the two `double` expressions will be evaluated right away to produce the resulting list `[4,8,12,16]` and

assign it to `r`. However, in lazily evaluated languages, evaluations are deferred until results are needed by other computations. For example, the evaluation will occur when the programmer wants to get the first element in the resulting quadrupled list. Specifically, the first `double` will request a result from the second "inner" `double`. Then, the second "inner" `double` will be executed to produce the doubled list. Finally the outer `double` takes the resulting doubled list as input, outputs the quadrupled list, and returns its first element. The biggest advantage of a lazy language is to improve code modularity without sacrificing code efficiency. In the field of programming languages, laziness is often referred to as the call-by-name evaluation scheme.

This isn't quite right—the elements are quadrupled one by one.

call-by-name is slightly different from "lazy"

~~redundant title~~ 2.4.2 Dependently-Typed Dependent Haskell

For a long time, Haskell researchers have striven to add dependent types in Haskell, as dependent types introduce much more precise expressions of program specifications. Although Haskell doesn't yet support dependent types in full, several extensions to its current type system have made exceptional progress.

In 2012, Eisenberg and Weirich introduced the `singletons` library that simulates dependently-typed programming in Haskell. Through dependent types, programmers can reason about programs through computations on the indexed types. For example, when concatenating two lists of booleans, each of length `m` and `n`, the output list can be specified to have type $\text{Vec}(\text{Plus } m \ n) \text{ Bool}$, indexed on the type variable `Plus m n` representing the sum of the two lengths of the original lists. However, Haskell enforces two separate scopes in programming functions – the terms executed at run time and the types checked at compile time. Functions in these two scopes cannot be used interchangeably, so singleton types are necessary to break the gap. Specifically, through the `promote` function provided in the `singletons` library, we can directly promote the `Plus` function defined in the term level to the type level using Template Haskell and apply it to reason about the dependent types. In addition, through the `genSingletons` function, the `singletons` library supports automatic singleton datatype definitions from any defined general datatype and provides a synonym for every singleton datatype generated. For example, `SNat n` is a synonym for the promoted singleton datatype of `Nat`, that can be used in the type level to retrieve the `n`th element from a list. Finally, the `singletons` function provided in the `singletons` library can generate the promoted type family for any function as well as automates a runtime version of the function that works with singleton types [5].

but you haven't introduced vectors yet

but it's much more than just a scoping issue!

not helpful to your readers.

include citations

Beyond Singletons, Gundry and Eisenberg both conducted research on dependently-typed Haskell using GADTs and type families (Detailed definitions and examples are provided in section 2.4 Program Verification). Eisenberg, extending Gundry's work, brings Haskell to a full-spectrum dependently typed language. Eisenberg's work, now implemented in the Glasgow Haskell Compiler (GHC) 8.0, finally blurs the distinction between type and kind, unifying expressions and types in Dependent Haskell. Unlike Gundry, who limits type promotion to exclude lambda expressions and unsaturated functions, Eisenberg enables promoting any expression available at run time to type [4]. To clarify the relationship between type and kind, consider the number 5, which is a term. Similar to any strongly typed languages such as Java, 5 has type `Int` in Dependent Haskell. Then, one can go one level up and classify all types, including `Int`, `Bool`, `Vec 5 Bool` etc, by kinds. Every normal data type that have values is classified by the special kind `Type` (or denoted as `*`), which is analogous to how every integer value is classified by the type `Int`.

In short, kinds can be understood as the type of types, which classify types, as how types classify terms. give an example of a type whose kind is not Type. ↗ when? in 8.0? or in my thesis?

Besides the unification of expressions and types, Eisenberg introduces the `'` operator to bring ordinary term-level functions directly into type-level, and the universal quantifier \forall to enrich the logic. Another major contribution of Eisenberg's work is pattern matching on dependent types, such as pattern matching on the indexed type `n` in the dependent type `Vec n a` (where `n` is the length of the vector and `a` is the type of its element).

Finally, Dependent Haskell continues as Haskell to be a partial language, without compile-time termination checking. Dependent Haskell only guarantees correctness of fully evaluated proofs, thus only allowing the well-typed values. For example, though erroneous proofs like `1 : ~ : 2`, a propositional equality GADT representing the logic $1 = 2$, might be verified, they will never be fully evaluated down to values. The erroneous logic `1 : ~ : 2` won't be evaluated to value since it doesn't follow the rules defined in Dependent Haskell's type system, therefore not well-typed.

this needs to
be much more
concrete to make sense.

2.5 Program Verification in Dependent Haskell

After a brief overview of Haskell and Dependent Haskell, this section focuses on Dependent Haskell's verification techniques. The section is structured in a tutorial manner, where concepts are introduced sequentially and gradually build up to two program verification examples as `insert` and `concat`.

2.5.1 Algebraic Data Types

Type system based program verification is all about types. Beyond the primitive types, `Int`, `Bool` etc, which are defined in the standard library, one can also define one's own types. Similar to the C structs, new types can be introduced through the `data` keyword.

For example, a list of elements of type `a` is defined recursively in Haskell as:

```
data List a = Nil | Cons a (List a)
    this is not from a library
```

The library list type is defined above with two value constructors, `Nil` and `Cons`. Each constructor specifies a value that the list type could have. Specifically, `Nil` is the base case representing the empty list and `Cons`, often written as `(:)`, is the inductive step that combines one more element to an existing list. The symbol `|` represents logical or, so the above code means that the `List` type whose elements are all of type `a` can have values either an empty list or a combination of an element of type `a` and another list of elements of type `a`. In functional programming and type theory, any type defined with value constructors is called an algebraic data type (ADT). ↗ ADTs are not unique to functional programming but this isn't a proof — it's a data structure!

Similarly, one can formally define the natural numbers recursively in Dependent Haskell as:

```
data Nat where
    Zero :: Nat
    Succ :: Nat -> Nat
    why use different
    syntax here?
```

where `Zero` represents the smallest natural number 0, and `Succ` represents the nonzero numbers as they are all increments of a smallest natural number.

2.5.2 Generalized Algebraic Data Types

It follows that, by using dependent types, one can encode the length of a list into its type as a natural number, by parameterizing the list by its length and its element type. Below we introduce the dependently typed vector definition in Dependent Haskell:

```
data Vec :: Nat -> Type -> Type where
    Nil :: Vec Zero a
    (:>) :: a -> Vec n a -> Vec (Succ n) a
infixr 5 :>
```

Instead of a simple listing of value constructors, the definition of `Vec` extends that of `List` by explicitly clarifying the type signature for each constructor, instantiating the algebraic data type (the why is this dependently typed? explain to your readers `Vec Zero a` and `Vec (Succ n) a`) as return values. Notice that both `List` and `Vec` resemble a linked list how will your readers know this? `? type`. Since functional languages don't modify any state of the input list, input list is often copied over to form the output list. In the above definition, the `:>` symbol represents the `Cons` operator as before that combines an element with another vector to produce a new vector. Also, `infixr 5 :>` defines the `Cons` operator to be right-associative, i.e. `1 :> 2 :> 3 :> Nil` will be evaluated as `1 :> (2 :> 3)` and results in the vector `[1, 2, 3]`. In general, an ADT whose value constructors have explicit type instantiations and match promoted term-level ADT patterns to refine type-level statements is called a Generalized Algebraic Data Type, or a GADT.

2.5.3 Type Class

Finally, we introduce the concept of a typeclass. Typeclasses can be understood as an interface that categorize a group of behaviors. It contains many types that all support and implement the behavior the associated typeclass describes. In Haskell, type classes provide a structured way to control ad hoc polymorphism, or commonly known as overloading. For example, `Ord` is a typeclass for all types that have an ordering, for example `Int`, `Double` etc. All elements of the type that have an ordering can be compared to check if one is greater than, equal, or less than another. To specify the usage of a certain type class, take `Ord` as an example, programmers can pass in the type variable `a` as in `Ord a` to specify that type `a` is part of the `Ord` type class and has an ordering. needs to be much more concrete `?`

2.5.4 A Simple Verification Example

With the introduced concepts as parameterized dependently-typed GADT `Vec` and the `Ord` typeclass of all types with orderings, we begin with a simple example on how program verification works in Haskell.

Specifically, one can define the `insert` function in Dependent Haskell to insert an element to a vector.

```
insert :: Ord a => a -> Vec n a -> Vec (Succ n) a
insert elm Nil      = elm :> Nil
insert elm (x :> xs) = if elm > x then x :> (insert elm xs) else elm :> x :> xs
```

Elements in the list of type `a` are all comparable. Since in the type signature `insert :: Ord a => a -> Vec n a -> Vec (Succ n) a`, we define `a` to belong in the `Ord` type class, as in `Ord a =>`. Also, the above implementation uses GADT pattern matching to

classify which case to apply. Specifically, `Nil` is to pattern match the original input list with the empty list, and `(x :> xs)` is to pattern match the original input list with at least an element. The above insert algorithm may be understood imperatively as follows: the element `elm` is inserted to the list if the list input is originally empty. Otherwise, as the list contains ordered elements, one may compare the value of `elm` with the smallest element `x` in the original nonempty list, and either recursively insert the new element into the original list or directly prepend the element to the front of the list. Notice that when the above `insert` algorithm is repeatedly applied, the eventual output list will be sorted.

Similar to writing test cases, program verifications also focus on the specification of each program. For `insert`, the length of the list is incremented by 1 every time one inserts a new element. Therefore, the return type for function `insert` can be defined as `Vec (Succ n) a`, encoding the length incrementation in the returned GADT.

This is the verification part—but it's not explained how this works.

2.5.5 Type Family

To perform some more involved computations in type level, such as addition or multiplication, Dependent Haskell introduces what's known as a type family. For example when concatenating two `Vecs`, one needs to verify the specification such that the resulting `Vec` has length the sum of the two concatenated `Vecs`.

Type family is a function defined at type level that returns a type. It gives us an extended ability to compute over types and is the core of type-level computation. For example, to add two type-level `Nats` together, we define the type family in Dependent Haskell as:

```
type family Plus (a :: Nat) (b :: Nat) :: Nat where
  Plus Zero b      = b
  Plus (Succ a') b = Succ (Plus a' b)
infixl 6 Plus
```

The implementation of `Plus` is exactly the same as the ordinary `(+)` implementation in the term level, except that `Plus` is one ~~match up~~ⁱⁿ the type level. Also, `infixl 6 +` defines the `(+)` operator to be left-associative, i.e. $1 + 2 + 3$ will be evaluated as $(1 + 2) + 3$ on default. Again, we apply pattern matching on the first summand of plus. If the first summand `a` is 0, then we define our base case such that any sum of 0 and another natural number `b` is just the other natural number `b`, or in short $0 + b = b$ for all natural number `b`. Otherwise, if the first summand is a non-zero natural number, represented as `Succ a` where `a'` is the natural number one less than `a`, then one recursively applies `Plus` on the smaller natural number `a'` and the other natural number `b` to obtain the sum and add one to obtain the result. Equivalently, in algebraic equations, if $a = 1 + a'$ then the operation is essentially $a + b = (1 + a') + b = 1 + (a' + b)$ from associativity of plus.

is it Plus or +?

2.5.6 Program Verification with Dependent Types and Type Families

Finally, applying both dependent types and type families, one can verify the operation of concatenating two lists. The operation `concat` can be recursively implemented and verified in Dependent Haskell as:

```

concat :: Vec n a -> Vec m a -> Vec (Plus n m) a
concat Nil v2      = v2
concat (x :> xs) v2 = x :> (concat xs v2)

```

*walk us through
how this type-checks*

The implementation is similar as before, applying pattern matching on the first parameter and recursion on the smaller list. After concatenating two lists, the resulting vector is expected to have length the sum of the lengths of the two original vectors, as specified in the dependent types `Vec (Plus n m) a` applying the type family `Plus` in the typing level.

2.6 FStar

`F*` is an ML-like functional programming language aimed at program verification [14]. It is developed at Microsoft Research, MSR India and Inria. `F*`'s syntax is similar to OCaml and F# in the ML family. ML is short for the Meta Language invented for proving theorems. In 2016, Swamy et al. introduced the current, completely redesigned version of `F*`. The new `F*` is a general-purpose, verification-oriented and effectful programming language. `F*` is said to be general-purpose as it is expressive enough to support programming in various practical domains. `F*` is also verification-oriented as it applies formal verification techniques to prove the correctness of programs written in `F*`. Finally, `F*` is also effectful, for it allows programs with side-effects, such as I/O, exceptions, stateful processing, or programs that potentially diverge.

*really?
proving theorems*

2.6.1 Introduction to `F*`

2.6.1.1 `F*` Basics

As a functional programming language, `F*` shares most concepts with Haskell, like statically ~~typ~~ typing, strictly typing, currying etc. There are only small differences between `F*`'s syntax and that of Dependent Haskell. Below is an example of `F*`'s implementation of factorial:

```

val factorial: x:int{x>=0} -> Tot int
let rec factorial n =
  if n = 0 then 1 else n * (factorial (n - 1))

```

In `F*`, programmers also specify the function signature before the actual implementation. All function signatures start with the keyword `val` and the type signature is introduced after a single colon instead of a double colon in Haskell. Also, all function implementations start with the keyword `let` and if the function is recursively called, the keyword `rec` will be specified between `let` and the function name.

`F*` supports both dependent types and refinement types. Instead of defining a new GADT for natural numbers, in order to restrict the input of `factorial` to natural numbers, one can also refine the type of `x` to `int{x >= 0}`. This type with logical predicate `x >= 0` is called a refinement type. Also notice that the function `factorial` doesn't just return an `int` but `Tot int`. The prefix `Tot` is a special derived form of the `PURE` effect, categorizing all pure functions (the concept of pure is defined in section 2.3.1). This `Tot` effect specifies the function `factorial` to be total, i.e. guaranteed to terminate. As an effectful language, `F*` categorizes the programming effects and in addition divergence in a lattice, where some effects are part of other effects. For example, the effect `PURE` on

I don't think this is the term you want here

total functions is a **sub-effect** of the divergent effect DIV, which is again a sub-effect of the exception effect EXN and finally the effect ALL categorizes every effect defined [14]. Specifically, Swamy et al. define these effects as monads that are used by F* to discharge verifications using both the Satisfiability Modulo Theories (SMT) solver (discussed in section 2.6.3) and manual proofs.

2.6.1.2 Simple Inductive Types

Similar to Dependent Haskell, F* allows creating new types using the `type` keyword. We show the definition of the standard library function `list` as follows:

```
type list 'a =
| Nil : list 'a
| Cons : hd:'a -> tl:list 'a -> list 'a
```

The definition of the `list` datatype in F* follows a similar syntax as Haskell. One can use `'a` to denote an arbitrary type (like generics in Java) and use `=` instead of `where` to start the actual type definition. Notice that the type constructors `Nil` and `Cons` both have to begin with a capital letter. *but there is pattern matching here*. The pattern matching is indicated by the `|` keyword and the constructor types follow by a single colon. Finally, F* requires all return types to annotate effects. In both type constructors for `list`, when effects aren't specified, the default effect `Tot` is automatically inferred by F* for the result `list 'a`.

With the `list` definition, the `length` function can be easily programmed to return the length of any input list. Specifically, one applies pattern matching on the input list `l` with syntax `match l with` and uses the `|` symbol to separate cases. The `length` function defined on `list` is widely used in the refinement logic to specify conditions for proving lemmas.

```
val length: list 'a -> Tot nat
let rec length l = match l with
| [] -> 0
| _ :: tl -> 1 + length tl
```

There are two new syntax introduced in the `length` definition, two syntactic sugar for the list constructors where `[]` represents `Nil` and `_ :: tl` represents `Cons _ tl`. In the first pattern matching case, one defines the length of an empty list, `Nil`, to be 0. In the second pattern matching case, one recursively defines the length of a list to be one more than the length of the sublist starting from the second element. The `_` symbol is commonly known as a wildcard, representing the value we don't care about when performing pattern matching.

Finally, the same `concat` function introduced in section 2.4.7 can be defined in F* as follows:

```
val concat : #a:Type -> list a -> list a -> Tot (list a)
let rec concat #a l1 l2 = match l1 with
| [] -> l2
| hd :: tl -> hd :: concat #a tl l2
```

In the above implementation, an extra argument `#a:Type` is introduced in the type definition. When preceded by the `#` symbol, the type `a` is turned into an implicit argument. Implicit arguments, always preceded by a `#`, tell the F* type system to try inferring

the type automatically instead of always requiring programmers to provide manually. It is a kind of type inference that ease the complexity of writing code in a strictly typed environment. Similarly to the '`a`' argument in the `list` definition, the implicit argument `#a` of type `Type` denotes a generic type.

2.6.2 Intrinsic and Extrinsic Style of Verification

`F*` in general supports two approach for verifying programs. Programmers can prove properties either by enriching the types of a function or by writing a separate lemma. These two ways towards verification are often called the intrinsic and extrinsic style.

2.6.2.1 Proving List Concatenation Intrinsically through Type Enrichment

In the `concat` example for concatenating two lists, one wants to prove the property that the length of the resulting list is equal to the sum of the lengths of the two input list. `F*` supports verifying the property by indexing on the length of the list using dependent types. With the definition of the `vec` GADT, an implementation of the proof is shown below, similar to the one presented in section 2.4.7 in Dependent Haskell.

```
type vec (a:Type) : nat -> Type =
| Nil : vec a 0
| Cons : hd:a -> n:nat -> tl:vec a n -> vec a (n + 1)
```

This definition of the dependently typed `vec` is analogous to the `Vec` GADT in Dependent Haskell as defined in section 2.4.3, except that we apply the type parameter `(a:Type)` before colon to bring the generic type `a` in scope for all type definitions below.

Then the verified `concat` function is exactly the same as in Dependent Haskell except the need to introduce two implicit arguments to bring the two list lengths in scope.

```
val concat : #a:Type -> #n:nat -> #m:nat -> l1:vector a n -> l2:vector a m ->
    Tot (vector a (n + m))
let rec concat #a #n #m l1 l2 = match l1 with
| Nil      -> l2
| Cons hd tl -> Cons hd (concat #a tl l2)
```

2.6.2.2 Proving List Concatenation Extrinsically through Lemmas

In addition to the intrinsic verification style, `F*` also supports extrinsic verifications through lemmas. Going back to the weak typing using `list` and with the predefined `length` function, one can prove the lemma as follow:

```
val concat_pf: l1:list 'a -> l2:list 'a
    -> Lemma (requires True)
        (ensures (length (concat l1 l2) = length l1 + length l2))
let rec concat_pf l1 l2 = match l1 with
| []       -> ()
| hd::tl -> concat_pf tl l2
```

`concat` is a total function that guarantees that it will terminate. Hence, it can be directly lifted wholesale to the logic level and be reasoned about in the postcondition

following the keyword `ensures`. In the `concat_pf` lemma, the property is valid in all conditions, so the precondition following the keyword `requires` is simply defined as `True`. With both precondition and postcondition specified, one can prove the lemma recursively. Notice that the `()` symbol introduce here is equivalent to `Refl` in Dependent Haskell, short for reflexivity, and denotes the satisfaction of the property.

2.6.3 Proof Automation

A lot of research has been done to reduce the tedious work through automating the proof processes. Automated theorem proving is a subfield of mathematical logic that proves mathematical theorems through computer programs. Satisfiability Modulo Theories (SMT) solvers provide a good foundation for highly automated programs [10].

Satisfiability Modulo Theories (SMT) problems are a set of decision problems that generalize boolean satisfiability (SAT) with arithmetic, fixed-sized bit-vectors, arrays, and other related first-order theories. Z3 is an efficient SMT Solver introduced by Microsoft Research that is now widely adopted in program verifications [3]. Dependent Haskell doesn't support proof automation at any level, but the F* language achieves semi-automation through a combination of SMT solving using Z3 and user-provided lemmas.

With refinement types, F* can discharge the proof on the `concat` example above automatically using Z3.

```
val concat : 11:list 'a -> 12:list 'a -> Tot (1:list 'a{length 1 = length 11 + length 12})
let rec concat 11 12 = match 11 with
| []          -> 12
| hd :: t1 -> hd :: concat t1 12
```

The return type for the list 1 can be specified using the refinement type as `list 'a{length 1 = length 11 + length 12}`. Then F* will generate verification conditions (VC) based on the refinement type and eventually pass to the Z3 SMT-solver for the verification.

2.6.4 Categorizing effects in F*

2.6.4.1 Dijkstra Monad and Hoare Logic

explain without using monads

F* categorizes programming effects (as introduced in Section 2.3) using what's known as a Dijkstra monad. Dijkstra monad is a special kind of monad that incorporates Hoare logic to classify effects into preconditions and postconditions. Hoare logic also introduces what's known as a weakest precondition. The term "weak" here refers to how general an assertion is, so weakest precondition is the most general precondition, i.e. logical predicate on the program inputs, needs to establish the postcondition. For a terminating division function on natural numbers, any predicate excluding zero in the divisor x is a valid precondition, such as `int{x > 5}` or `int{x > 10}`. Yet the weakest precondition can be uniquely defined as `int{x > 0}`.

Swamy et al. observes that in the context of dependently typed programming, the weakest precondition function can be modeled as a monad [14]. We won't go into details of the Dijkstra monads but the core idea is to categorize the weakest preconditions for each

effect introduced before, such as PURE or DIV, as monads. Besides the complicated theories behind the categorization of programming effects, F* provides the `require-ensure` syntactic sugar to help with actual programming. F* programmers only need to provide the function postconditions following the `requires` keyword and preconditions following the `ensures` keyword, and F* can desugar the syntax to discharge actual proofs.

Through the effect system designed using Dijkstra monads, F* provides a powerful mechanism, called monadic returns, to enrich function types extrinsically. Specifically, for any total function `e` that can be proved to terminate, both its type and definition can be lifted into the logic and reason about it either in the subtyping predicates of refinement types or in the postconditions following keyword `ensures`.

2.6.4.2 Proving Termination

As we are mostly interested in how F* handles potentially divergent functions, we will focus on the `Tot` and `DIV` effect and introduce F*'s approach towards termination checking.

F* by default requires termination checking for every function, since divergent functions may generate unsound verification conditions (VCs) and eventually end up proving some erroneous logic. F* introduces a new termination criterion based on a well-founded partial order \prec over all terms. A partial order is a binary relation on a set that is reflexive, antisymmetric and transitive. For example, the \leq ordering on integers is a partial order, since for every integer `a`, `b` and `c`, it follows that `a ≤ a` (reflexive), `a ≤ b` but `b ≥ a` (antisymmetric), and finally if `a ≤ b` and `b ≤ c` then `a ≤ c` (transitive). F* insists each iteration of a recursive function be applied on a strictly smaller domain, such that the parameter of the function is guaranteed to decrease in the defined partial order during each execution. The F* type-checker provides four built-in well-founded ordering that can be automatically applied in proving termination, but it also accepts decreasing metrics explicitly provided by the programmer following the `decreases` keyword. Below we summarize these four default metrics and provide an example in each:

(1) The natural number ordering: For instance, the natural number parameter in the `factorial` example in section 2.6.1.1 is decreasing in every iteration.

(2) For any function `f: x:t → Tot t'`, the curried function (See section 2.1 for the discussion on currying) `f v` for some `v:t` is "less than" the fully defined function `f`. For instance, consider the `length` function introduced in section 2.6.1.2, we would have `length (1 :: 2) ≺ length`. This metric on currying is very rarely used, so we won't go into details. Instead of applying this metric on currying, the termination proof on `length` is through default metric (3) discussed next.

(3) The sub-terms of an inductively defined term is less than the term itself: For instance, in the `concat` example in section 2.6.2.1 and 2.6.3, every recursively call on the tail of the input list (sublist starting from the second element) is smaller than the original call on input list since the length of the tail is shorter than that of the input list.

(4) The lexicographical ordering: By default, F* checks the decreasing metric on each non-function-typed argument in its order of application. For instance, in the `concat` example in section 2.6.2.1 and 2.6.3, if the order of `11` and `12` is switched but the function

definition still pattern matches on 11, then F* will first check the decreasing metric on the first parameter 12. Since 12 stays the same, the termination checker continues to check the second parameter 11 whose length indeed decreases by default metric (3). Hence the `concat` function with parameters switched is still proved total.

When the function is proved total, one can specify it using the `Tot` effect, notifying the F* effect system that the function is guaranteed terminating. If, instead, one can't find any termination metric or the function indeed diverges, F* provides the `DIV` effect to categorize all these potentially divergent functions. Once a function is marked as `DIV`, F* will turn off all termination checking on that function.

3 Approach and Novelty

To gain a better understanding of the language design decisions and verification techniques, We first carefully studied the previous literature on each language. In general, our assessments of Dependent Haskell and F* are based on implementing and proving a wide spectrum of algorithms in both languages. These programs are written so that we can have a much more in-depth experience with both languages and explore their features more thoroughly.

We start from implementing the datatype of a length-indexed vector in both Dependent Haskell and F*. Then we first verify various functions that support types of finite length vectors from Haskell's `Data.List` module, including `length`, `head`, `tail`, `init`, `last`, `snoc`, `reverse`, `and`, `or`, `null` etc. We then focus on the verification for various sorting algorithms, such as insertion sort, mergesort and quicksort. Finally, we examine the capability of proving divergent functions by surveying through two examples. The first example is the proof of the Peano division property $(a * b) / b = a$ using the divergent definition of division with zero divisors. The second proof is the equivalence of multi-step and big-step evaluators in simply typed lambda calculus.

We expect to structure the comparison mainly focusing on two aspects: user interface and approaches towards potentially non-terminating functions. As far as we know, we are the first researchers to directly compare the two languages, Dependent Haskell and F*.

User interface refers to the interaction between programmer and the programs, specifically programming simplicity when implementing programs in Dependent Haskell or F*. We are interested in examining how these two language designs differ in respect to the ease of use measured from syntactic verbosity and library/documentation resources. As both languages are verification oriented, in addition to general programming syntax, we will also include a comparison on proof complexities for each verification technique. We believe that a good programming language should provide a clear, concise and user-friendly interface with a relatively detailed library and documentation supports.

The approach towards potentially non-terminating functions is measured by the "strength" and completeness of the type systems. A type system is said to be "strong" if every property proved using it is guaranteed valid in logic. Conversely, a type system is said to be complete if every logically valid property with respect to the system's semantics can be proved in the system. Both type systems in Dependent Haskell and F* are "strong" and complete when we restrict the range of functions to those that are provably terminating, called total functions. However, when we extend the system to consider potentially non-

terminating functions, either functions that diverge or are provably hard to prove convergence, ensuring both strength and completeness in the type system is a huge challenge. Without compile time termination checking, Dependent Haskell seems to be complete and is able to prove properties on logically valid properties for divergent functions, yet its type system is not "strong" enough, as it verifies some logically erroneous proofs and requires run time checking to ensure the program correctness (Research in progress). On the other hand, F* is provably "strong" as one can ensure the termination of all total functions of the Tot effect. Yet it's unknown whether F*'s type system is complete with its categorization of potential divergent functions using the DIV effect.

4 Current Results

As summarized in Section 3, we have implemented various functions that support types of finite length vectors from Haskell's `Data.List` module, three machine-checked sorting algorithms – Insertion Sort, MergeSort and QuickSort, and two proofs on properties from potentially divergent functions – Peano division with a divergent division definition and the equivalence of multi-step and big-step evaluator in Simply Typed Lambda Calculus.

4.1 User Interface: General Language Design and Support

4.1.1 Syntactic Verbosity

In terms of general language design and support, we conclude that Dependent Haskell is syntactically verbose while F* has a fairly friendly user interface.

4.1.1.1 Dependent Haskell

Dependent Haskell enforces phase separation between the term and type level. Thus, to apply a function defined in the term level to a dependent type in the type level, programmers are often expected to repeat implementations of the same logic.

The first type of duplication is in data type definitions. Dependent Haskell enables dependently typed programming through the use of singletons. Singleton is a special type, also known as a singleton type, that bridges the gap between run-time values and compile-time types. Each singleton type is indexed by a type variable of the promoted kind (see section 2.4.2 on definition of kind). For example, `SNat :: Nat -> Type`, a GADT indexed by a type of kind `Nat` and returns a new type `Type`(or `*`), is a singleton type for the `Nat` datatype on natural numbers. A singleton type is a type with exactly one value, where each type variable indexing the singleton type is uniquely mapped to the term of that type. Hence there is an isomorphism between singleton types and the associated values, enabling term-level computation and type-level computation to go hand in hand [5].

Take the natural numbers as an example, there are three different definitions – the original datatype `Nat`, the kind `'Nat`, and the derived singleton type `SNat`. In 2012, Yorgey et al. introduced the datatype promotion mechanism that automatically promotes datatypes up a level to kinds and the data constructors to type-level data [18]. Hence, the kind definition `'Nat` can be automatically derived from the original datatype, reducing programming redundancy down by a level. However, prior to the introduction of the `singletons` library, programmers still need to manually implement the singleton type,

even though the definition of a singleton type is a straightforward extension to that of the original datatype.

```
-- original datatype Nat definition
data Nat :: * where
  Zero :: Nat
  Succ :: Nat -> Nat

-- singleton type SNat definition
data SNat :: 'Nat -> * where
  SZero :: SNat 'Zero
  SSucc :: SNat n -> SNat ('Succ n)
```

As shown above, `'Zero` and `'Succ` are both automatically promoted type-level data from the `Zero` and `Succ` data constructors. Also, the `SZero` and `SSucc` data constructors are indexed on the promoted `'Nat` type and help specify the dependency between the compile-time type parameters (`'Zero`, `'Succ` `'Zero` etc) and the run-time term values (`Zero`, `Succ` `Zero` etc).

The `singletons` library introduced by Eisenberg and Weirich removes the above code duplication. Using Template Haskell, the library can automatically generate the definition for singleton types. However, the programming overhead still remains, as the library only frees programmers from the tedious singleton definitions but doesn't hide the nature of the internal singletons encoding. Programmers are still expected to have a deep understanding of the phase separation and the singleton types, and to take full charges of the function implementations using the library-generated definitions.

The second type of duplication comes from function definitions. Functions such as adding two natural numbers, need to be specified once in the term level for the original datatypes, once in the type level using type families, and once in the term level for the singleton types using the type-level type family definition.

```
-- term-level function plus definition
plus :: Nat -> Nat -> Nat
plus Zero n = n
plus (Succ m) n = Succ (plus m n)

-- type-level type family Plus definition
type family Plus (m :: Nat) (n :: Nat) :: Nat where
  Plus Zero n = n
  Plus (Succ m') n = Succ (Plus m' n)

-- term-level function sPlus definition for singleton types
sPlus :: SNat m -> SNat n -> SNat (Plus m n)
sPlus SZero n = n
sPlus (SSucc m) n = SSucc (sPlus m n)
```

As shown above, we define the term-level `plus` function, adding two natural numbers and obtaining the result of their sum. The base case is defined, in mathematical terms, as $0 + n = n$ and the inductive case as $(1 + m) + n = 1 + (m + n)$. Similarly, following the exact same logic, we define the type family `Plus` for adding two natural numbers in the type level. That is to say that instead of using the `plus` function in the actual implementation, the `Plus` type family can be applied to the type signature as a quantifier

over the dependent types. Finally, we implement the function `sPlus` on the singleton types. The only difference between the definition of `sPlus` and that of `plus` is that the operation is quantified over the type parameter of the singleton type instead of directly performing on the terms. Specifically, pattern matching on the singleton type `SNat m` gives us insights into both the type variable of kind `Nat` and the associated term-level variable of type `Nat`.

The above three versions of the same addition algorithm on natural numbers add in lots of additional burdens when using the language. To free programmers from these repetitive and tedious works, the `singletons` library also automates the function definition process. Through Template Haskell, the library supports promoting term-level functions and directly reuse them in the type level. It also supports enriching functions with richer types, extending the term-level definitions to directly apply on singleton types.

With the `singletons` library, we only need to provide the original data type definition for `Nat` and the term-level function definition for `plus` as shown below:

```
$(singletons [d|
  data Nat = Zero | Succ Nat
  deriving Eq

  plus :: Nat -> Nat -> Nat
  plus Zero    n = n
  plus (Succ m) n = Succ (plus m n)
  []])
```

Notice that the above code follows the Template Haskell syntax, where it starts with the `$` symbol and follows by the normal Haskell datatype and function definitions enclosed in the quoting syntax `[d| ... |]`. This `$()` syntax is known as a Template Haskell splice, whose contents are evaluated at compile time [5]. Specifically, in the natural numbers addition example, the `singletons` function takes in the Template Haskell quote containing an abstraction of both the datatype definition on `Nat` and the function definition on `plus`. It then generates a list of Template Haskell declarations including all five definitions – `SNat`, `Plus`, `sPlus` – introduced before. Finally, these Template Haskell declarations are inserted back into Dependent Haskell and are defined to use in all program verifications.

The `singletons` library indeed reduces programmer's burden to manually implement all functions at different level through the automation of singleton data type definition, and the generation of type family and enriched functions on singleton types. Yet, the library still doesn't resolve the syntactic verbosity in Dependent Haskell. The `singletons` library only helps lessen the effort users need to spend on the duplicated declarations, but it doesn't abstract away any detail of the actual definitions. As a result, programmers still need to fully understand how type family and singleton work to compile their programs and to complete the verifications. Therefore, even with the assistance of both the datatype promotion technique and the `singletons` library, Dependent Haskell's syntax remains a little clunky and is not very friendly to users new to programming verification.

4.1.1.2 F*

On the other hand, F* is syntactically concise. It supports both type-level specifications through dependent types and type-level computations through refinement types.

For example, the type signature for the `concat` function introduced in section 2.4.7 can be specified using dependent type as `val concat: #a:Type -> #n:nat -> #m:nat -> Vec n a -> Vec m a -> Vec (n + m) a`. F*'s approach towards dependent types is similar to that of Dependent Haskell, except that F* frees programmers from the repetitive type family implementation of `Plus`. Through the `T-Ret` typing rule using the idea from monadic returns (we won't go into detail for the typing rule, but a brief introduction is provided in section 2.6.4.1), F* can automatically enrich the types of a total function and reflect the existing term-level implementations into logic for reasoning. Monadic returns in F* are similar to the `return` typing rule in monads, returning the weakest precondition categorized as a monad after proving the post-condition on it. Specifically in the example of `concat`, the library function `(+)` is automatically reflected to the type level and is applied in the dependent type `Vec (n + m) a`.

The same verification for concatenating two lists can be specified using refinement type, as `val concat: #a:Type -> #n:nat -> #m:nat -> 11:Vec a -> 12:Vec a -> r:Vec a{len r = len 11 + len 12}`. Instead of indexing the length of the vector to form dependent types, we can specify the intuitive condition `len r = len 11 + len 12` as a logical predicate to the resulting type using subtyping.

4.1.2 Library and Documentation Supports

Dependent Haskell, extending Haskell with dependent types, has been officially implemented in GHC 8.0 and is now a part of Haskell. As Haskell is a mature, industrial-strength programming language, Dependent Haskell inherits all its resources. Specifically, Haskell provides a central package archive called Hackage where packages are introduced, maintained, and supported. In addition, Hoogle is a holistic search engine for Haskell APIs which provides advanced queries through function names such as `sum`, and even through approximate type signatures such as `int -> int -> int`. As Haskell becomes more popular in the real world, there emerges a plethora of online tutorials and textbooks with detailed explanations of syntax, concepts and algorithms. Since the introduction of Dependent Haskell in 2014, tutorials, such as schoolofhaskell, have gradually incorporated Dependent Haskell into their documentations. Official resources for Dependent Haskell are also available on Richard Eisenberg's Github repository for his PHD dissertation (<https://github.com/goldfirere/thesis>).

However, as a completely redesigned language that is relatively new, F* is still in lack of many resources. Besides various conference publications, F* has an official website organizing all of its resources (<https://www.fstar-lang.org>). Most of the documentations are through its interactive tutorial (<https://www.fstar-lang.org/tutorial/>) and examples in the official FStar Github repository (<https://github.com/FStarLang/FStar>). In addition, source code for each library function can be found in the `ulib` folder of its Github repository (<https://github.com/FStarLang/FStar/tree/master/ulib>).

4.1.3 Proof Complexity

Regarding verification techniques, Dependent Haskell doesn't support any proof interaction or automation. Unlike most verifiers, it chooses to avoid compile-time termination checking but supports, through monads, most general programming effects like IO and Error etc. Verification in Dependent Haskell is guaranteed both sound and complete for total functions, but could get really tedious even for relatively simple sorting algorithms, like insertion sort, mergesort or quicksort.

(details on DH complexity on insertion sort, mergesort or quicksort are to be filled)

F^* , on the other hand, provides a flexible combination of SMT-based proof automation and manually constructed proofs. In addition, F^* can use SMT solver, such as Z3, to match SMT patterns provided in the manually provided lemmas to assist with user-constructed proofs. F^* supports theorem proving through either enriching function types (officially referred as the intrinsic style) or by writing separate lemmas on properties (officially referred as the extrinsic style) [14]. It categorizes effects into Tot, DIV etc using Dijkstra monads, requires explicit annotation of the effect in type signatures and enforces termination checking on all total functions.

4.2 Approaches to Potentially Non-terminating Functions

To find out Dependent Haskell and F^* 's different verification techniques towards non-terminating functions, we start from exploring the Peano division example (whose complete implementations are provided in Appendix). Peano naturals are a simple way to encode natural numbers with the `Zero` base value and the `Succ` recursive successor, each time incrementing the value by one. For example, natural number 0 is expressed as `Zero`, 1 is expressed as `Succ Zero` and 2 is expressed as `Succ (Succ Zero)` etc. Using the Peano division example, we want to prove the property that the quotient of a product of two natural numbers a, b and b is equal to a , i.e. $a * b / b = a$. As an educational example to simulate non-termination, we allow zero divisors in our definition of `division` and try to prove the property when the divisors are indeed positive, using the divergent division definition.

4.2.1 Dependent Haskell might be Complete towards Potentially Divergent Functions

Dependent Haskell doesn't require termination checking. With the help from type families and singletons, Dependent Haskell can successfully verify the Peano division property.

We first define the GADT `Peano` and the type level operations `+`, `-`, `*`, `/` in type families as follows:

```
data Peano where
  Zero :: Peano
  Succ :: Peano -> Peano

type family (a :: Peano) + (b :: Peano) :: Peano where ...
type family (m :: Peano) - (n :: Peano) :: Peano where ...
type family (a :: Peano) * (b :: Peano) :: Peano where ...
type family (a :: Peano) / (b :: Peano) :: Peano where ...
```

These four operations are used in the type signature of `peanoDivisionPf` defined below to prove the Peano division property.

Then we define the singleton type `SPeano` for natural numbers and specify the type signatures for the term level `plus`, `minus` and `mult` that are used in the actual implementation of the proof.

```

data SPeano n where
  SZero :: SPeano Zero
  SSucc :: SPeano n -> SPeano (Succ n)

  plus :: SPeano m -> SPeano n -> SPeano (m + n)
  minus :: SPeano m -> SPeano n -> SPeano (m - n)
  mult :: SPeano m -> SPeano n -> SPeano (m * n)

```

Finally, we prove the property $(m * n) / n = m$ for positive n in `peanoDivisionPf` using the three lemmas `plusZero`, `plusSucc` and `plusMinus`. We can define, using singletons, the positive divisor as `SPeano (Succ n)`, so the property becomes $(m * (Succ n)) / (Succ n) = m$. The `peanoDivisionPf` on the property and the three lemmas applied are shown below:

```

plusZero :: SPeano m -> m + Zero :~: m
plusZero SZero      = Refl
plusZero (SSucc m) =
  case (plusZero m) of Refl
-> Refl

plusSucc :: SPeano m -> SPeano n -> m + (Succ n) :~: Succ (m + n)
plusSucc SZero _      = Refl
plusSucc (SSucc m) n =
  case (plusSucc m n) of Refl
-> Refl

plusMinus :: SPeano m -> SPeano n -> (m + n) - n :~: m
plusMinus m SZero      =
  case (plusZero m) of Refl
-> Refl
plusMinus m (SSucc n) =
  case (plusSucc m n) of Refl
-> case (plusMinus m n) of Refl
-> Refl

peanoDivisionPf :: SPeano m -> SPeano (Succ n) ->
  (m * (Succ n)) / (Succ n) :~: m
peanoDivisionPf SZero _          = Refl
peanoDivisionPf (SSucc m) succ_n@(SSucc n) =
  case plusMinus (mult m succ_n) succ_n of Refl
-> case plusSucc (mult m succ_n) n of Refl
-> case peanoDivisionPf m succ_n of Refl
-> Refl

```

To understand the `peanoDivisionPf` proof, let's first discuss the three lemmas applied. The lemma `plusZero` proves that `Zero` is an additive identity of the Peano natural numbers. By definition of additive identity, the sum of every Peano natural number and the `Zero` Peano natural number is the Peano natural number itself, i.e. $m + \text{Zero} :~: m$ for some Peano natural number m . The lemma `plusSucc` proves that $m + (\text{Succ } n) :~: \text{Succ } (m + n)$, or in algebraic form, $m + (1 + n) = 1 + (m + n)$ for some Peano natural numbers m and n . Notice that `:~:` is a propositional GADT similar to the equal

sign that denotes the logical equality in the type level. Finally, the lemma `plusMinus` proves the property on minus, where $(m + n) - n : \sim: m$.

Dependent Haskell provides the `@` syntax to assist with pattern matching. In `succ_n@(SSucc n)`, `(SSucc n)` is called an invisible argument, prefixing with `@`. `(SSucc n)` exposes `n` for function definition and brings it to scope. Following `@`, `succ_n` is the explicit value that stands for `SSucc n`. That is to say that every `(SSucc n)` in the following code can be expressed as `succ_n` [4]. Assume the inductive hypothesis that $(m * (Succ n)) / (Succ n) : : m$, the inductive step of the proof is shown step by step as follows:

```
(Succ m * Succ n) / Succ n
-- by definition of (*)
= (m * Succ n + Succ n) / Succ n
-- by plusSucc
= Succ (m * Succ n + n) / Succ n
-- by definition of (/)
= Succ Zero + (Succ (m * Succ n + n) - Succ n) / Succ n
-- by definition of (+)
= Succ Zero + ((Succ (m * Succ n) + n) - Succ n) / Succ n
-- by definition of (+)
= Succ (((Succ (m * Succ n) + n) - Succ n) / Succ n)
-- by plusSucc (same as above, opposite direction)
= Succ (((m * Succ n) + Succ n - Succ n) / Succ n)
-- by plusMinus
= Succ ((m * Succ n) / Succ n)
-- by inductive hypothesis
= Succ m
```

In summary, Dependent Haskell doesn't require any level of termination checking at compile time, so divergent functions such as division by zero can still be reasoned easily and used to prove the Peano division property valid in logic. We still need more examples to determine whether Dependent Haskell is indeed "strong".

4.2.2 F* is Incomplete towards Potentially Divergent Functions

F*, based on the incomplete Hoare Logic, doesn't support extrinsic proofs on non-terminating functions. In Hoare Logic, the transition from one program state to another is encoded in Hoare triples. Hoare triple is of the form $P \ C \ Q$, where P and Q are called assertions in the predicate logic and denote the precondition and postcondition correspondingly. C is the command that is being executed. Hoare logic is a mature way to categorize programming states, but unfortunately it is incomplete, as it cannot derive the logically valid specifications if the execution C cannot be proved to terminate. As a result, F* allows proving divergent programs intrinsically, instead through type enrichment, as shown in the example of Peano Division.

First, we define the `peano` GADT in F* as follows:

```
type peano: Type =
| Zero : peano
| Succ : peano -> peano
```

Then we show the type signatures of `plus`, `minus`, `mult` and `division` on natural numbers as follows:

```

val maximal: int -> int -> Tot int
let maximal a b = if a > b then a else b

val plus: a:peano -> b:peano -> Tot (c:peano{toNat c >= toNat a && toNat c >=
toNat b})
val minus: a:peano -> b:peano -> Tot (c:peano{toNat c = maximal (toNat a -
toNat b) 0})
val mult: a:peano -> b:peano -> Tot (c:peano)
val division: a:peano -> b:peano -> Tot (c:peano)

```

Notice that `toNat` is a total function that takes in a `peano` GADT and returns the corresponding natural number, i.e. `toNat Zero` = 0 and `toNat (Succ Zero)` = 1 etc. Our verified `plus` function ensures that the sum is greater than or equal to both summands, and `minus` ensures that the difference is either 0 or equivalent to the difference of the two corresponding natural numbers.

For division, as we include the possibility that the divisor of `division` could be zero, we must mark the effect for `division` as `DIV`, meaning divergent. When we try to verify the property $(a*b)/b = a$ following the same logic as in Dependent Haskell, we find ourselves in trouble. That's because the `division` function of effect `DIV` cannot be directly reflected to verification conditions (VC) through F^* 's monadic return typing rule. Hence, we cannot use the divergent `division` function in `ensures`, i.e. the postcondition of the lemma.

In this educational Peano division example, we can easily modify the type signature to define a terminating `div_terminating` function. Thus, one workaround to prove the lemma $(a*b)/b = a$ is by explicitly establishing the logical connection between the divergent function `division` and the total function `div_terminating`. It follows that we can complete the proof by applying `div_terminating` to reason about the property $(a*b)/b = a$ in the postcondition of the lemma.

Below we show the functions `div_terminating` and `division` as well as the type signature for the proof `peanoDivPf`:

```

val div_terminating: a:peano -> b:peano -> Pure (c:peano)
  (requires toNat b > 0)
  (ensures fun y -> True)
  (decreases (toNat a))
let rec div_terminating a b =
  if a = Zero then Zero
  else let denom = (div_terminating (minus a b) b)
    in plus denom (Succ Zero)

val division: a:peano -> b:peano -> Div (c:peano)
  (requires True)
  (ensures (fun y -> toNat b > 0 ==> y == div_terminating a b))
let rec division a b =
  if a = Zero then Zero
  else let denom = (division (minus a b) b)
    in plus denom (Succ Zero)

val peanoDivPf: a:peano -> b:peano -> Lemma
  (requires (toNat b > 0))
  (ensures (div_terminating (mult a b) b == a))

```

Specifically, the total function `div_terminating` takes in two peano numbers – dividend `a` and divisor `b` – and returns the quotient `c`. Using the syntactic sugar `requires` and `ensures` as introduced in section 2.6.2.2, we specify the precondition to allow only nonzero divisors and specify the postcondition to be universally true for any inputs (fun `y` -> `True` is a lambda expression as in the familiar form `\y -> True` that is beyond the scope of the current discussion). To prove the totality of `div_terminating`, that is to prove that `div_terminating` indeed terminates, we specify our decreasing metric as `toNat a`. The metric `toNat a` is always decreasing since when we recursively call `div_terminating` on `(minus a b)` and `b`, we are in the `else` case where `a` is a nonzero natural number and by definition of `minus`, we always have `(minus a b) < a`. Therefore, with the decreasing metric, the function `div_terminating` is guaranteed to terminate by the well-ordering principle on natural numbers in algebra.

It follows that in the divergent function `division`, we can ensure that for any precondition, as long as the divisor for `division` is nonzero, the result of `division` is equivalent to the result of `div_terminating`.

We call the proof `peanoDivPf` partial since the property only holds for a certain precondition. Also, though the introduction of the intermediary total function `div_terminating` completes the proof, this approach inevitably requires programmers to know the exact reason of non-termination. Therefore we believe that F^* is not generally applicable, therefore incomplete, in verifying the potentially divergent functions in most practical programs.

5 Related Work and Contributions

5.1 Language Comparison

There are many programming languages in use today and new ones are created every year. Various work has been done to compare programming languages. Through comparison, researchers find advantages and limitations in each and evaluate to propose insights for future research [6].

In the realm of programming languages, there are general-purposed programming languages like Pascal, Haskell and C/C++, interactive proof assistants like Coq, Agda and Isabelle/HOL, and SMT-based semi-automated program verification tools like Dafny, Vampire [8] and WhyML. Wiedijk compared a list of proof assistants on mathematical theorems, focusing on their strength of logic and level of automation [15] [16]. Feuer and Gehani edited papers comparing and assessing general purposed languages Ada, C and Pascal [6]. Filiâtre compared Vampire and Dafny in proving steps of the progress proof [13] [7].

(to be filled: summarize what they did and found)

As the needs grow for general purpose yet verification oriented programming languages, four functional programming languages have been recently invented, namely Dependent Haskell, Liquid Haskell, F* and Idris. As an ongoing field of research, to my best knowledge, these four communities generally conduct research separately and have rarely explored each other's approaches in-depth.

At the 45th ACM POPL Student Research Competition, Xu and Zhang presented their collaborated work on the comparison among three programming verification techniques in Dependent Haskell, Liquid Haskell and F* [17]. As an exploratory research, they found promising aspects in each language and proposed more detailed directions for further comparisons. In this thesis, we aim to provide an in-depth comparison between Dependent Haskell and F*, with a focus on their different verification approaches towards non-terminating functions. Through in-depth assessments of the verification techniques in both languages, we aim to find the more desirable approach and to recommend future directions for research in termination checking.

5.2 Dependent Haskell

to be filled

5.2.1 Functional Dependency

to be filled

5.2.2 Constrained Type Families

to be filled

5.2.3 Promoting Functions to Type Families in Haskell

to be filled

5.3 F*

to be filled

5.3.1 Relative Completeness in Hoare Logic

to be filled

5.3.2 Dijkstra Monads

to be filled

5.3.3 Hereditary Substitutions

to be filled

5.4 The Zombie Language

to be filled (Details about how the Zombie Language attempts to verify divergent functions)

Appendices

.1 Insertion Sort

.1.1 Dependent Haskell

to be filled

.1.2 F*

to be filled

.2 Peano Division

.2.1 Dependent Haskell

```
{-# LANGUAGE GADTs, TypeInType, ScopedTypeVariables, StandaloneDeriving,
   TypeFamilies, TypeOperators, AllowAmbiguousTypes, TypeApplications,
   UndecidableInstances, DataKinds #-}

module PeanoDivisionDH where

import Data.Kind ( Type )
import Data.Type.Equality ( (=:~:)(..) )
import Prelude hiding ( div )

data Peano where
  Zero :: Peano
```

Second Draft

```

Succ :: Peano -> Peano
deriving (Eq, Ord, Show)

type family (a :: Peano) + (b :: Peano) :: Peano where
  Zero + b      = b
  (Succ a) + b = Succ (a + b)
infix 6 +

type family (m :: Peano) - (n :: Peano) :: Peano where
  Zero - _       = Zero
  m - Zero      = m
  (Succ m) - (Succ n) = m - n
infix 6 -

type family (a :: Peano) * (b :: Peano) :: Peano where
  Zero * _       = Zero
  (Succ a) * b = a * b + b
infix 7 *

type family (a :: Peano) / (b :: Peano) :: Peano where
  Zero / _       = Zero
  m / n         = (Succ Zero) + (m - n) / n
infix 7 /

data SNat n where
  SZero :: SNat Zero
  SSucc :: SNat n -> SNat (Succ n)
deriving instance Show (SNat n)

plus :: SNat m -> SNat n -> SNat (m + n)
plus SZero n      = n
plus (SSucc m) n = SSucc (plus m n)

minus :: SNat m -> SNat n -> SNat (m - n)
minus SZero _     = SZero
minus m SZero     = m
minus (SSucc m) (SSucc n) = minus m n

mult :: SNat m -> SNat n -> SNat (m * n)
mult SZero _      = SZero
mult (SSucc m) n = plus (mult m n) n

plusZero :: SNat m -> m + Zero :~: m
plusZero SZero    = Refl
plusZero (SSucc m) = case (plusZero m) of Refl -> Refl

plusSucc :: SNat m -> SNat n -> m + (Succ n) :~: Succ (m + n)
plusSucc SZero _ = Refl
plusSucc (SSucc m) n = case (plusSucc m n) of Refl -> Refl

--  (m + (Succ n)) - (Succ n)

```

```
-- = Succ (m + n) - (Succ n), by plusSucc
-- = (m + n) - n, by definition of (-)
-- = m, by induction hypothesis
plusMinus :: SNat m -> SNat n -> (m + n) - n :~: m
plusMinus m SZero      = case (plusZero m) of Refl -> Refl
plusMinus m (SSucc n) =
  case (plusSucc m n) of Refl
-> case (plusMinus m n) of Refl
-> Refl

-- (Succ m * Succ n) / Succ n
-- = (m * Succ n + Succ n) / Succ n, by definition of (*)
-- = Succ (m * Succ n + n) / Succ n, by plusSucc
-- = Succ Zero + (Succ (m * Succ n + n) - Succ n) / Succ n, by definition of (/)
-- = Succ Zero + ((Succ (m * Succ n) + n) - Succ n) / Succ n, by definition of (+)
-- = Succ (((Succ (m * Succ n) + n) - Succ n) / Succ n), by definition of (+)
-- = Succ (((m * Succ n) + Succ n - Succ n) / Succ n), by plusSucc (same as
  above, opposite direction)
-- = Succ ((m * Succ n) / Succ n), by plusMinus
-- = Succ m, by induction hypothesis
peanoDivisionPf :: SNat m -> SNat (Succ n) -> (m * (Succ n)) / (Succ n) :~: m
peanoDivisionPf SZero _           = Refl
peanoDivisionPf (SSucc m) succ_n@(SSucc n) =
  case plusMinus (mult m succ_n) succ_n of Refl
-> case plusSucc (mult m succ_n) n of Refl
-> case peanoDivisionPf m succ_n of Refl
-> Refl
```

.2.2 F*

```
module PeanoDivision

type peano: Type =
| Zero : peano
| Succ : peano -> peano

val toNat: peano -> Tot nat
let rec toNat a = match a with
| Zero    -> 0
| Succ a' -> 1 + toNat a'

val max: int -> int -> Tot int
let max a b = if a > b then a else b

val plus: a:peano -> b:peano ->
  Tot (c:peano{toNat c >= toNat a && toNat c >= toNat b})
let rec plus a b = match a with
| Zero    -> b
| Succ a' -> Succ (plus a' b)
```

Second Draft

```

val minus: a:peano -> b:peano ->
  Tot (c:peano{toNat c = max (toNat a - toNat b) 0})
let rec minus a b = match a with
  | Zero    -> Zero
  | Succ a' -> match b with
    | Zero    -> Succ a'
    | Succ b' -> minus a' b'

val mult: a:peano -> b:peano -> Tot (c:peano)
let rec mult a b = match a with
  | Zero    -> Zero
  | Succ a' -> plus (mult a' b) b

val div: a:peano -> b:peano{toNat b > 0} -> Tot (c:peano)
(decreases (toNat a))
let rec div a b =
  if a = Zero then Zero
  else plus (div (minus a b) b) (Succ Zero)

val plusZero: m:peano -> Lemma (ensures plus m Zero = m)
let rec plusZero m = match m with
  | Zero    -> ()
  | Succ m -> plusZero m

val plusSucc: m:peano -> n:peano -> Lemma (ensures plus m (Succ n) = Succ (plus
  m n))
let rec plusSucc m n = match m with
  | Zero    -> ()
  | Succ m -> plusSucc m n

val plusMinus: m:peano -> n:peano -> Lemma
(requires True)
(ensures minus (plus m n) n = m)
let rec plusMinus m n = match n with
  | Zero    -> plusZero m
  | Succ n -> match plusSucc m n with () -> plusMinus m n
  (*
    minus (plus a (Succ b')) (Succ b')
    = minus (Succ (plus a b')) (Succ b') by plusSucc a b'
    = minus (plus a b') b'
    = a by minusPlus a b'
  *)

```



```

val peanoDivPf: m:peano -> n:peano{toNat n > 0} -> Lemma
(requires True)
(ensures (div (mult m n) n == m))
let rec peanoDivPf m n = match m with
  | Zero    -> ()
  (*
    div (mult Zero b) b = div Zero b = Zero
  *)

```

```

| Succ m ->
  match plusMinus (mult m n) n with ()
-> match peanoDivPf m n with
-> match plusSucc m Zero with
-> plusZero m
(*
  div (mult (Succ a') b) b
= div (plus (mult a' b) b) b
= plus (div (minus (plus (mult a' b) b) b) b) (Succ Zero)
= plus (div (mult a' b) b) (Succ Zero) by plusMinus (mult a' b) b
= plus a' (Succ Zero) by peanoDivPf a' b
= Succ (plus a' Zero) by plusSucc a' Zero
= Succ a' by plusZero a'
*)

```

References

- [1] Stephen Cass. The 2017 top programming languages, 2017. *where is this published?*
- [2] Coq development team. *The Coq proof assistant reference manual*. LogiCal Project, 2004. Version 8.0. *Capitalize*
- [3] Leonardo De Moura and Nikolaj Bjørner. Z3: An efficient *smt* solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*, pages 337–340. Springer, 2008.
- [4] Richard A. Eisenberg. *Dependent Types in Haskell: Theory and Practice*. PhD thesis, University of Pennsylvania, 2016.
- [5] Richard A. Eisenberg and Stephanie Weirich. Dependently typed programming with singletons. In *ACM SIGPLAN Haskell Symposium*, 2012.
- [6] Narain Gehani. *Comparing and assessing programming languages: Ada, C, and Pascal*. Prentice Hall, 1984. *?*
- [7] Sylvia Grewe, Sebastian Erdweg, and Mira Mezini. Automating proof steps of progress proofs: Comparing vampire and dafny. In *Vampire@IJCAR*, pages 33–45, 2016. *Capitalize*
- [8] Laura Kovács and Andrei Voronkov. First-order theorem proving and vampire. In *International Conference on Computer Aided Verification*, pages 1–35. Springer, 2013.
- [9] K. Rustan M. Leino. Dafny: An automatic program verifier for functional correctness. In Edmund M. Clarke and Andrei Voronkov, editors, *Logic for Programming, Artificial Intelligence, and Reasoning*, LPAR’16, pages 348–370. Springer Berlin Heidelberg, 2010.
- [10] K. Rustan M. Leino. Automating induction with an *smt* solver. In *International Workshop on Verification, Model Checking, and Abstract Interpretation*, pages 315–331. Springer, 2012.

- ~~use the conference citation instead in skel~~
- [11] James McKinna. Why dependent types matter. In *ACM Sigplan Notices*, volume 41, pages 1–11. ACM, 2006.
 - [12] Ulf Norell. *Towards a practical programming language based on dependent type theory*. PhD thesis, Department of Computer Science and Engineering, Chalmers University of Technology, SE-412 96 Göteborg, Sweden, September 2007.
 - [13] Benjamin C Pierce. *Types and programming languages*. MIT press, 2002.
 - [14] Nikhil Swamy, Cătălin Hrițcu, Chantal Keller, Aseem Rastogi, Antoine Delignat-Lavaud, Simon Forest, Karthikeyan Bhargavan, Cédric Fournet, Pierre-Yves Strub, Markulf Kohlweiss, Jean-Karim Zinzindohoue, and Santiago Zanella-Béguelin. Dependent types and multi-monadic effects in F^* . In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '16. ACM, 2016.
 - [15] Freek Wiedijk. Comparing mathematical provers. In *International Conference on Mathematical Knowledge Management*, pages 188–202. Springer, 2003.
 - [16] Freek Wiedijk. *The seventeen provers of the world: Foreword by Dana S. Scott*, LNCs volume 3600. Springer, 2006.
 - [17] Rachel Xu and Xinyue Zhang. Comparisons among three programming verification techniques in dependent haskell, liquid haskell and f^* . In *Proceedings of the 45th Annual ACM SIGPLAN Symposium on Principles of Programming Languages*, POPL '18. ACM, 2018. *not quite. This is where POPL papers go!*
 - [18] Brent A Yorgey, Stephanie Weirich, Julien Cretin, Simon Peyton Jones, Dimitrios Vytiniotis, and José Pedro Magalhães. Giving haskell a promotion. In *Proceedings of the 8th ACM SIGPLAN workshop on Types in language design and implementation*, pages 53–66. ACM, 2012.